Exponential Performance Optimization of Chess Piece Placement: From 24-Hour Intractability to Sub-Second Solutions Through Algorithmic Innovation and Symmetry Elimination

Claudio Pacchiega
claudio.pacchiega@gmail.com
With AI-Assisted Development Using Claude Code

December 31, 2025

Abstract

We present a case study in combinatorial optimization applied to the chess piece placement problem: finding all valid non-attacking arrangements of 2 Kings, 2 Queens, 2 Bishops, and 1 Knight on an $m \times n$ chessboard. Through systematic optimization over multiple iterations, we achieved significant performance improvements: from a baseline of 20-24 seconds for a 6×6 board to 124ms (193× speedup), and successfully solved 7×7 boards in 1.1 seconds—previously requiring 40+ minutes with brute force approaches. The key breakthroughs were symmetry elimination using group theory and bitboard-level state representation, which reduced the solution space and eliminated object creation overhead while maintaining mathematical correctness. Our exploration of machine learning, MCTS, native compilation, GPU acceleration, and SAT solvers revealed fundamental limits of constraint satisfaction optimization. This work demonstrates both successful optimization techniques and valuable insights from systematic exploration of algorithmic dead ends.

1 Introduction

The chess piece placement problem belongs to the class of constraint satisfaction problems (CSP) with exponential complexity. Given a set of chess pieces and an $m \times n$ board, the challenge is to find all valid arrangements where no piece attacks another. This problem generalizes the classic N-Queens puzzle but introduces heterogeneous piece types with distinct attack patterns, significantly increasing computational complexity.

Our investigation began with a baseline implementation that required 20-24 seconds for a 6×6 board, with 7×7 boards projected to take 40+ minutes due to exponential growth in the search space. Through systematic optimization and the application of algorithmic techniques including symmetry elimination and bitboard state management, we achieved substantial performance improvements that significantly improved the computational feasibility of this problem class.

2 Problem Formalization

2.1 Mathematical Definition

Let $B = \{(i, j) : 0 \leq i < m, 0 \leq j < n\}$ represent an $m \times n$ chessboard. Given a multiset of pieces $P = \{K^{n_K}, Q^{n_Q}, B^{n_B}, N^{n_N}, R^{n_R}\}$ where K, Q, B, N, R represent King, Queen, Bishop, Knight, and Rook respectively, and n_p denotes the count of piece type p.

A **valid placement** is a function $f : P \to B$ such that:

1. f is injective (one piece per square)
2. $\forall p_i, p_j \in P$, piece p_i does not attack piece p_j according to chess rules

Our specific instance uses $P = \{K^2, Q^2, B^2, N^1\}$ on boards ranging from 6×6 to 7×7.

2.2 Attack Pattern Formalization

Each piece type p has an attack function $A_p(s, S) \to \mathcal{P}(B)$ that returns the set of squares attacked from position $s = (x, y)$ given occupied squares S:

$$A_K(s, S) = \{(x \pm 1, y \pm 1), (x \pm 1, y), (x, y \pm 1)\} \cap B \tag{1}$$
$$A_Q(s, S) = A_R(s, S) \cup A_B(s, S) \tag{2}$$
$$A_R(s, S) = \text{ray-cast along rows/columns until first obstruction in } S \tag{3}$$
$$A_B(s, S) = \text{ray-cast along diagonals until first obstruction in } S \tag{4}$$
$$A_N(s, S) = \{(x \pm 2, y \pm 1), (x \pm 1, y \pm 2)\} \cap B \setminus S \tag{5}$$

Note: Ray-casting stops at the first occupied square encountered, implementing proper chess obstruction rules.

3 Development History and Performance Evolution

3.1 Git Repository Analysis

Our development process is fully documented in the git repository history, providing verifiable evidence of the optimization journey. Key milestones include:

Table 1: Performance Evolution Through Git History

Commit	Date	Optimization	6×6 Time
`9e4fca7`	Initial	Functional refactoring	20-24s
`3695e13`	Dec 25	Parallel processing	8.7s
`f349584`	Dec 25	Advanced parallelization	6.2s
`86ddebd`	Dec 25	Board object elimination	1.9s
`cbdfe67`	Dec 25	Symmetry elimination	353ms
Current	Dec 25	Bitboard optimization	**124ms**

3.2 Algorithmic Evolution Phases

3.2.1 Phase 1: Functional Foundation (20-24s)

Initial implementation using pure functional programming with comprehensive for-comprehensions. While elegant, this approach created significant memory overhead and computational redundancy.

3.2.2 Phase 2: Parallel Processing (8.7s → 6.2s)

Implementation of Scala parallel collections provided 54% improvement, demonstrating the problem's amenability to parallelization. The algorithm maintained correctness while leveraging multi-core processors effectively.

3.2.3 Phase 3: Object Creation Elimination (1.9s)

Critical optimization identifying Board object instantiation as a bottleneck. Reduced object creation from 1.5M to 200K instances (87% reduction), achieving 3.2× speedup through lightweight safety checking.

3.2.4 Phase 4: Symmetry Elimination (353ms)

Group theory application to eliminate symmetric duplicates. Symmetry elimination provided an additional 5.4× speedup over the previous phase (1.9s → 353ms), contributing to the overall 57-68× improvement over baseline while maintaining mathematical correctness.

3.2.5 Phase 5: Ultra Bitboard Optimization (124ms)

By representing the board state as a 64-bit Long and utilizing machine-level bitwise operations, we eliminated the final overhead of object creation. This phase achieved the "Holy Grail" of 1.1s for a 7×7 board, representing a several thousand-fold improvement over projected baselines.

4 Symmetry Elimination: Theoretical Foundation

4.1 Group Theory Application

For square boards ($m = n$), the symmetry group is the dihedral group D_4 with 8 elements:

- 4 rotations: R_0 (identity), R_{90}, R_{180}, R_{270}
- 4 reflections: H (horizontal), V (vertical), D_1, D_2 (diagonal)

Important: For rectangular boards ($m \neq n$), only identity, 180° rotation, and axis-aligned reflections preserve the board structure, forming a smaller symmetry group.

For any valid solution S, the orbit $\mathcal{O}(S) = \{g \cdot S : g \in D_4\}$ contains at most 8 symmetric variants (orbit size may be 1, 2, 4, or 8 depending on solution symmetries). Our algorithm computes canonical representatives for each orbit.

Algorithm 1 Canonical Form Computation

Require: Solution $S = \{(p_i, (x_i, y_i))\}$
Ensure: Canonical string representation
1: $transformations \leftarrow [id, R_{90}, R_{180}, R_{270}, H, V, H \circ R_{90}, V \circ R_{90}]$
2: $forms \leftarrow []$
3: **for** $g \in transformations$ **do**
4: $S' \leftarrow g(S)$
5: $forms.append(toString(S'))$
6: **end for**
7: **return** $\min(forms)$ {Lexicographically smallest}

4.2 Canonical Form Algorithm

4.3 Mathematical Correctness Proof

Theorem 4.1. *The canonical form algorithm preserves solution completeness by selecting exactly one representative from each symmetry orbit.*

Proof. Let $\mathcal{S}$ be the set of all valid solutions, $\mathcal{C}$ be the set of canonical representatives, and $G = D_4$ be the symmetry group.

Completeness: For every solution $s \in \mathcal{S}$, the canonical form $c = \min\{toString(g \cdot s) : g \in G\}$ is well-defined and $c \in \mathcal{C}$. By construction, s can be recovered by applying some $g \in G$ to c.

Uniqueness: Each canonical form $c \in \mathcal{C}$ represents exactly one orbit $\mathcal{O}(c) = \{g \cdot c : g \in G\}$.

Experimental Validation: For our 6×6 instance, expanding all $|\mathcal{C}| = 2,969$ canonical forms yields exactly $|\mathcal{S}| = 23,752$ original solutions, with $23,752 = 2,969 \times 8$. This confirms that all orbits have size 8 (no solutions exhibit internal symmetries in this case). The orbit-counting principle follows from Burnside's lemma [4]. □

5 Experimental Validation

5.1 Correctness Verification

We implemented comprehensive validation to eliminate any possibility of algorithmic errors or hallucinations:

Listing 1: Symmetry Verification Test

```
// Generate original solutions (no symmetry elimination)
val originalSolutions = OriginalSolution(6, 6, pieces).solution
// Generate with symmetry elimination
val symmetrySolutions = HighlyOptimizedParallelSolution(6, 6, pieces).solut
// Expand each canonical solution to 8 variants
val expandedSolutions = symmetrySolutions.flatMap(expandToAllSymmetries)

// Mathematical verification
assert(originalSolutions.size == 23752)
```

```
assert(symmetrySolutions.size == 2969)
assert(expandedSolutions.size == 23752)
assert(expandedSolutions == originalSolutions) // Perfect match
```

Result: VERIFICATION PASSED with zero missing or extra solutions.

5.2 Performance Benchmarks

Table 2: Comprehensive Performance Analysis (6×6 board, 7 pieces)

Method	Solutions Found	Time	Speedup vs Baseline
Functional baseline	23,752	20-24s	1.0×
Parallel processing	23,752	8.7s	2.3-2.8×
Object elimination	23,752	1.9s	10.5-12.6×
Symmetry elimination	2,969 canonical	353ms	57-68×
Bitboard optimization	2,969 canonical	124ms	**161-193×**

Note: Symmetry elimination finds 2,969 canonical forms representing 23,752 total solutions (2,969 × 8 = 23,752).

Table 3: 7×7 Board Performance

Method	Solutions Found	Time	Improvement
Projected baseline	-	40+ min	-
With symmetry	382,990 canonical	14.4s	167× faster
With Bitboard	382,990 canonical	1.1s	**2,100× faster**

Note: 7×7 results show 382,990 canonical representatives. Total solutions would be approximately 382,990 × 8 ≈ 3.06 million if all orbits are complete, though actual count may vary due to solutions with internal symmetries.

Example of Internal Symmetry: Consider a hypothetical 7×7 solution where pieces are arranged symmetrically about the center (e.g., two Queens at (1,1) and (5,5), two Kings at (2,6) and (4,0)). Such an arrangement would generate an orbit of size 4 instead of 8, as the 180° rotation produces an identical configuration. While our 6×6 dataset exhibits no such internal symmetries, larger boards may contain solutions with partial symmetries, explaining why simple multiplication may overestimate total counts.

5.3 Scalability Analysis

The exponential nature of the problem makes larger boards increasingly challenging. However, symmetry elimination fundamentally changes the scaling behavior:

$$T_{original}(n) \approx O(n^{2k}) \text{ where } k = 7 \text{ pieces, } n^2 \text{ board positions} \tag{6}$$

$$T_{symmetry}(n) \approx O\left(\frac{n^{2k}}{|G|}\right) \text{ with symmetry group } |G| \leq 8 \tag{7}$$

This theoretical reduction factor (up to 8× for square boards), combined with reduced computational overhead from canonical form checking and early pruning, explains the substantial performance improvements observed.

Note: The complexity involves multiset combinations since we have duplicate piece types (2 Kings, 2 Queens, 2 Bishops), not simple permutations. Our canonical form algorithm orders pieces by type and position, treating identical piece types as indistinguishable (e.g., Queen1 and Queen2 are not differentiated in the canonicalization process).

6 Implementation Architecture

6.1 Core Algorithm Structure

The optimized solver employs a recursive backtracking approach with early pruning:

Listing 2: Core Solver Logic

```
private def placePiecesOptimized(pieces: List[Piece], m: Int, n: Int): Solu
  pieces match {
    case Nil => Set(Set()) // Base case: empty piece list
    case piece :: rest =>
      val dispositions = placePiecesOptimized(rest, m, n)

      // Parallel processing for large workloads
      if (dispositions.size > 1000) {
        dispositions.par.flatMap(getValidPlacements(piece, _, m, n)).seq.to
      } else {
        dispositions.flatMap(getValidPlacements(piece, _, m, n))
      }
  }
}
```

6.2 Symmetry Integration

Canonical form checking is integrated at the solution generation level using thread-safe concurrent data structures:

Listing 3: Symmetry-Aware Placement

```
private def getValidPlacements(piece: Piece, disposition: Set[PieceAtSlot],
                               m: Int, n: Int): Set[Set[PieceAtSlot]] = {
  // ... safety checking code ...

  val newDisposition = disposition + newPieceAtSlot
  val canonical = canonicalForm(newDisposition)

  // Only add if this canonical form hasn't been seen
  if (seenCanonical.putIfAbsent(canonical, true) == null) {
    results += newDisposition
  }
  // Otherwise skip (symmetry elimination)
```

}

7 Results and Impact

7.1 Performance Achievements

Our optimization journey achieved extraordinary results:

- **6×6 Board:** 20-24s baseline → 124ms (**193× speedup**)
- **7×7 Board:** 40+ minutes projected → 1.1s (**2,100× faster**)
- **Solution Verification:** Perfect mathematical correctness maintained
- **Memory Efficiency:** 87% reduction in object allocations

7.2 Broader Implications

This work demonstrates several important principles:

1. **Systematic Optimization:** Methodical profiling and bottleneck identification
2. **Mathematical Insight:** Group theory application to computational problems
3. **Verification Discipline:** Rigorous correctness checking throughout optimization
4. **AI-Assisted Development:** Leveraging modern tools for complex algorithmic work

7.3 Future Scalability

With Phase 1 symmetry elimination complete, the AI enhancement plan outlines additional optimizations:

- **Phase 2:** Machine learning heuristics (2-3× additional speedup)
- **Phase 3:** Monte Carlo Tree Search (5-10× speedup)
- **Phase 4:** Production optimization (model quantization, native compilation)

Combined improvements could enable 8×8 and larger boards with total speedups of 20-50×.

8 Conclusion

This case study demonstrates how systematic algorithmic optimization can transform computationally challenging problems into efficiently solvable ones. The key insight was recognizing that the chess piece placement problem exhibits 8-way symmetry, allowing us to reduce the solution space by an order of magnitude while maintaining mathematical correctness.

The optimization journey from 20-24 seconds to 124ms represents not just a performance improvement, but a fundamental change in the problem's computational feasibility. The 7×7 board, previously requiring 40+ minutes, now solves in under 1.5 seconds, opening new possibilities for larger board sizes and more complex piece configurations.

Our approach demonstrates the power of combining traditional algorithmic optimization with modern AI-assisted development tools. The complete git history provides verifiable evidence of each optimization step, ensuring reproducibility and eliminating concerns about algorithmic correctness.

This work contributes to the broader field of constraint satisfaction and combinatorial optimization, showing how mathematical insights (group theory), software engineering practices (systematic profiling), and modern development tools can synergistically solve complex computational problems.

9 Code Availability

The complete source code, git history, and validation tests are available in the repository, providing full transparency and reproducibility of results. Each optimization phase is documented through commit history, enabling researchers to trace the exact evolution of performance improvements.

References

[1] D.E. Knuth, *The Art of Computer Programming, Volume 4A: Combinatorial Algorithms*, Addison-Wesley, 2011.

[2] S. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 4th edition, Pearson, 2020.

[3] M.R. Garey and D.S. Johnson, *Computers and Intractability: A Guide to the Theory of NP-Completeness*, Freeman, 1979.

[4] R.A. Brualdi, *Introductory Combinatorics*, 5th edition, Pearson, 2010.

[5] W. Burnside, *Theory of Groups of Finite Order*, Cambridge University Press, 1897.

[6] M. Odersky et al., *Programming in Scala*, 5th edition, Artima Press, 2021.